



Wakeel وكيل

The Agentic AI Platform for Government

Every government employee builds, tests and deploys AI agents — in plain English or Arabic, with no technical skills. Built for Abu Dhabi entities on a sovereign, open-source core.

Prepared for

Government of Abu Dhabi — entity leadership

Wakeel · Confidential — for discussion purposes

Date

July 2026

Status

Working prototype — live demo available

THE PROBLEM

Government runs on repetitive knowledge work — and AI adoption is bottlenecked by the few who can build

60–80%

of service-desk effort is triage, checking completeness, and drafting standard responses — highly automatable knowledge work¹

1:100s

ratio of AI-capable engineers to employees with automatable workflows inside a typical entity — the build capacity gap

Months

typical lead time from an idea ("auto-check permit applications") to a working solution through classic IT procurement

What entities experience today

- ✗ Every AI use case becomes an IT project with vendors and specifications
- ✗ Pilots stall: no way for business owners to test quality or trust outputs
- ✗ Citizen-data residency concerns block public-cloud AI tools
- ✗ Arabic is an afterthought in most AI tooling

What the UAE AI agenda demands


- ✓ AI in every government service — not in a central lab only
- ✓ Employees empowered to automate their own processes
- ✓ Sovereign deployment options and auditability
- ✓ Arabic and English as equal first-class languages

¹ Illustrative, based on published shared-services and service-desk benchmarks; to be validated per entity during pilot.

THE SOLUTION

Wakeel: describe the task — get a working, tested, deployed AI agent

One platform, one journey, designed for a non-technical government employee. Live today as a working prototype.



KEY DIFFERENTIATOR

Trust is the product: measured accuracy and self-healing agents

The reason government AI pilots stall is not building — it is trusting. Wakeel makes quality measurable and repairable by the business owner, not a data scientist.

The Feedback loop (live in the prototype)

- ✓ **Auto-generated test suites** — AI reads the agent's steps and writes realistic cases, including Arabic inputs and privacy edge cases
- ✓ **Employee feedback #1** — edit or replace any generated case; add your own
- ✓ **AI judge** — every run scored pass/fail with a reason; accuracy ring per agent
- ✓ **Employee feedback #2** — written rules the agent must obey ("always end with a bilingual signature")
- ✓ **Self-healing** — failures + feedback are sent back to the AI, which revises the agent and re-tests automatically

Measured on the live prototype

8 s

full test run incl. AI-judge verdict

11 s


self-heal cycle: revise agent + save

18 s

6 auto-generated test cases

"Was 67% accurate — self-healed to 100% on 9 July." The audit trail every entity can show its leadership.

Thin sovereign experience layer on a proven open-source engine



Single URL, single sign-on

One login controls the portal, embedded editor and APIs.
No tab-hopping, no second credentials.

Every action audited

Builds, tests, self-heals, publishes and provider changes are logged per user with timestamps.

No lock-in

Open-source engine, exportable agent packages (DSL), swappable models — the entity owns everything.

50 ready workflows across 10 Abu Dhabi entities — shipped in the marketplace today

ENTITY	EXAMPLE WORKFLOWS IN THE MARKETPLACE	ENTITY	EXAMPLE WORKFLOWS IN THE MARKETPLACE
Abu Dhabi Chamber	Membership renewal assistant · Certificate-of-origin checker · Trade inquiry responder	Education (ADEK)	School registration assistant · Scholarship reviewer · Teacher-license checker
Economic Development	Trade-license reviewer · Trade-name validator · Inspection report summarizer	Judicial Department	Case-filing assistant · Legal document summarizer · Contract clause extractor
TAMM / Digital Gov	Service-request triage · Complaint router · Customer sentiment analyzer	Abu Dhabi Police	Incident report classifier · Traffic-appeal reviewer · Community tip triage
Municipalities & Transport	Building-permit pre-check · Parking-fine appeal reviewer · Contractor document verifier	Federal Tax Authority	VAT registration screener · Penalty appeal drafter · Invoice compliance checker
Department of Health	Medical-license screener · Patient complaint classifier · Insurance pre-approval drafter	Identity & Citizenship	Visa pre-screener · Golden-visa eligibility checker · Family sponsorship assistant

Each installs as a live, editable agent in ~15–60 seconds and can be adapted to the entity's own rules, then tested and self-healed before go-live. The library grows as champions publish their own workflows back to the marketplace.

Trade-license renewal: from 3 manual touches to 1 approval

Before


Officer opens each renewal → checks tenancy contract, Emirates ID, outstanding fines → drafts reminder or approval letter manually → escalates unclear cases. ~12–18 min per application, quality varies by officer.

With Wakeel (built by the licensing team itself)

Agent: reads the application → verifies required fields & documents list → drafts a bilingual approval letter or a missing-items request → routes exceptions to the officer with a summary. Officer approves in one click.

Illustrative impact per 10,000 renewals / year¹

Officer time per application



~2,000 h

officer hours redirected to complex cases annually

< 1 day

from idea to tested agent — built in Wakeel, self-healed to target accuracy

The same pattern repeats across permits, appeals, screenings and inquiries — the marketplace makes each solved case a template for the next entity.

¹ Illustrative sizing for discussion; to be validated with entity volumetrics during the pilot.

Built for government constraints, not adapted to them

Data residency

- ✓ Per-entity deployment in Azure UAE North, sovereign clouds (G42 / Moro / Khazna) or on-premises
- ✓ Azure OpenAI UAE region keeps prompts and citizen data in-country
- ✓ Option for fully local models where classification requires it

Isolation & control

- ✓ One instance per entity — no shared tenancy, clean accreditation boundary
- ✓ Entities bring their own model API keys; switch providers in Settings without re-architecture
- ✓ Exportable agent packages — no vendor lock-in

Accountability

- ✓ Per-user activity log: who built, tested, healed, published — with timestamps
- ✓ Accuracy evidence per agent before go-live (test suites & scores)
- ✓ Roadmap: UAE Pass / Entra ID SSO, role-based approvals inbox

Licensing posture (validated against the Dify open-source license)

Commercial use as a backend is expressly permitted. Per-entity single-workspace deployments avoid the multi-tenant clause entirely; one branding remedy (or an enterprise white-label license from LangGenius) clears the console-branding clause. Full compliance brief available as a separate document.

Innoventures Champions: the community that makes Wakeel stick

Technology alone does not change how a ministry works — people do. Wakeel embeds the AI Champions programme (champions.innoventures.ae) directly in the platform as its Community tab.

1 • Champions network

Trained AI champions inside each entity — the first builders, the local help desk, and the quality gate for what their department publishes to the marketplace.

2 • Use-case exchange

Champions publish proven agents with their test suites and accuracy scores. ICP's golden-visa checker becomes tomorrow's template for another entity's eligibility screening.

3 • Learning hub & meetups

Structured learning paths, entity leaderboards and regular meetups — recognition and goals ("My Goals", "View Reports") keep momentum after the pilot ends.

The flywheel

Champion builds → tests & publishes → marketplace grows → next entity installs in minutes → new champions emerge. Each turn lowers the cost of the next use case and raises the quality bar.

Programme support offered with the platform

Champion onboarding & certification · entity playbooks · office hours · quarterly agent audits · executive reporting on adoption, accuracy and hours redirected.

Per-entity platform subscription — aligned with sovereignty and the license

Pilot

Proof of value — 6 weeks

- ✓ 1 entity, 1 cloud instance (Azure UAE North)
- ✓ 3 champion-built use cases + 5 marketplace installs
- ✓ Accuracy targets agreed up-front; go/no-go evidence pack

Fixed fee — infrastructure at cost

Scale

Entity subscription — annual

- ✓ Dedicated sovereign instance, unlimited employees & agents
- ✓ Marketplace + Champions programme access
- ✓ Support SLAs, upgrades, quarterly agent audits

Tiered by entity size; models billed on entity's own keys

Enterprise

Government-wide programme

- ✓ Multi-entity rollout with central champions office
- ✓ White-label license (LangGenius) folded in
- ✓ Shared marketplace governance & analytics

Structured with procurement

Weeks, not quarters

time-to-value per use case, with evidence of accuracy at go-live

OPEX-light

open-source core: no per-seat AI-platform license burn; entities pay for value & support

Compounding

every published workflow reduces the next entity's cost — the marketplace is the moat

From working prototype to government-grade programme

Now — live

- ✓ Chat-to-build agents (EN/AR), embedded editor, chat assistant
- ✓ Auto test generation, AI-judge accuracy, self-healing
- ✓ Marketplace (50 workflows), Community embed, activity audit
- ✓ Own API keys & provider switching, mobile UI

Pilot phase

- ✓ Knowledge-grounded agents (RAG on entity documents)
- ✓ Saved test suites & accuracy history per agent
- ✓ UAE Pass / Entra ID single sign-on
- ✓ Branding remedy / LangGenius enterprise license


Scale phase

- ✓ Scheduled & event-triggered agents; approvals inbox
- ✓ Role-based access & per-department spaces
- ✓ Cross-entity marketplace governance & analytics
- ✓ Sovereign local-model option for classified workloads

Principle: ship what employees can use safely today; add the governance layer as adoption grows — never the other way round.

THE ASK

A 6-week pilot with one entity — measured on evidence, not promises



What we need from the entity

One sponsor, two champions, one process with real volume, and access to sample (non-classified) documents.

What you get in week 6

A working agent portfolio your own team built — plus the audited numbers to take to leadership.



Wakeel وکیل

The Agentic AI Platform for Government · with Innoventures Champions

Wakeel · Confidential — for discussion purposes